


Executive summary

Our generative engine optimization audit of Nike evaluated a representative sample of **40 pages**, resulting in a weak overall score of **55.0 out of 100**. The most critical finding is that Nike’s AI visibility is severely restricted because several key pages are marked with noindex tags, which completely blocks AI crawlers from indexing and citing the brand's content. This technical barrier directly impacts Nike’s bottom line by ceding valuable real estate in AI-driven search results to competitors, ultimately reducing organic traffic and brand authority. To reclaim this search share, Nike must immediately **update crawler access rules in robots.txt to permit AI search engines to read public pages**. Additionally, the brand should **expand Organization structured data with verified sameAs profile links** to solidify brand identity, and **optimize content with passage-level answerability patterns** like direct definitions and FAQs that AI engines can easily extract and cite.

Score overview

Each row shows the current score on a 0–100 scale and a sparkline of past audits for this URL in the same output folder (e.g. sibling folders under audit_output/).



GA4 — AI traffic

Traffic pulled from GA4. If you have specified an AI channel in the setup wizard, this has been used to identify AI traffic. Otherwise, we use the source / medium dimension.

Connect GA4 for AI traffic data

When you run a crawl with `--ga4-property` (or set `GA4_PROPERTY_ID`), this section shows AI channel sessions, revenue share, and source/medium gaps beside your audit.

Prioritized action plan

At most **five** items per horizon (15 total cap). Overflow from Quick wins rolls into Medium-term, then Strategic. The **Strategic** column is oriented to long-cycle **content development**, **brand authority**, and **citability** (research-led publishing, governance, corroboration)—not quick technical patches. **Estimated lift** lines use conservative bands for the composite GEO score (category weights 40% AI visibility / 30% Technical / 30% Content)—they are **not** additive; validate with a follow-up crawl.


Quick wins	Medium-term	Strategic
0–30 days · highest return	30–90 days · high impact	90+ days · content, brand authority & citability programmes
1 Update the site's crawler access rules in robots.txt to ensure search engines and AI agents can access and index public product and landing pages. Est. +4–9 pts overall	1 Establish custom analytics tracking to measure AI-search referral traffic, conversion rates, engaged sessions, and query-to-page performance. Est. +2–6 pts overall	1 Strengthen Nike's authoritative digital footprint by securing high-quality external citations and optimizing key brand profiles like Wikipedia. Est. +2–6 pts overall
2 Audit and remove accidental 'noindex' tags on key templates to resolve the score cap and ensure critical pages are eligible for AI engine citation. Est. +4–10 pts overall	2 Conduct a comprehensive page speed and mobile usability audit to ensure slow load times are not hindering crawler access or user engagement. Est. +4–9 pts overall	2 Launch a research-led content program incorporating first-party data, expert commentary, and unique athletic insights to maximize information gain. Est. +4–10 pts overall
3 Enhance Organization structured data by adding verified social profile links ('sameAs') and extending JSON-LD schemas across priority product templates. Est. +3–8 pts overall	3 Verify and update key third-party brand sources, including Google Business Profile, major directories, and external reviews, to align brand signals. Est. +2–6 pts overall	3 Optimize content layouts for passage-level answerability by embedding clear definitions, structured tables, and direct answers under key headings. Est. +3–8 pts overall
4 Publish a short AI guide file (llms.txt) at the root directory to direct LLM crawlers to Nike's most important product, service, and policy pages. Est. +4–9 pts overall	4 Implement dateModified and datePublished metadata on time-sensitive content templates to signal freshness and relevance to AI crawlers. Est. +4–9 pts overall	4 Develop a comprehensive query-mapping framework to match common AI query shapes, such as product comparisons and troubleshooting, to dedicated pages. Est. +2–6 pts overall
5 Review canonical tags across primary templates to ensure search engines consolidate authority onto correct, listable URLs. Est. +5–12 pts overall		5 Establish an ongoing content governance framework to maintain source transparency, author disclosures, and scheduled freshness updates across all hubs. Est. +2–6 pts overall

Policy notes

- Separate decisions about training or dataset crawlers—not the same as AI search visibility fixes.
- Formulate an organizational policy on whether to allow or block ByteDance's Bytespider and other AI training crawlers based on data-use preferences.

Projected performance

Nike currently holds a composite score of **55.0**, indicating a weak overall AI search visibility posture. Resolving immediate crawler access and indexation blockers in the first phase is projected to elevate performance into the **65–75** range. Fully executing the medium-term tracking and strategic content optimization programs will position the brand to target the **80+** band, establishing strong authority across AI search engines.



Competitor comparison

Compare overall AI visibility, technical setup, and content quality side by side. Expand a peer row for structured notes: **where the competitor appears stronger** in the automated sample (with confidence and actions), and **where your site appears ahead**. Green items use validated crawl URLs only (HTTP 200, not soft-404 or noindex). Your row has no expandable notes.

SITE	OVERALL	AI VISIBILITY	TECHNICAL	CONTENT
https://nike.com	55.0	54.9	81.5	59.5
https://www.adidas.com	33.3	30.6	42.9	27.2
▸ What https://www.adidas.com is doing differently...				
https://www.underarmour.com	92.6	90.2	95.1	93.4
▸ What https://www.underarmour.com is doing differently...				
https://www.puma.com	85.1	80.7	88.2	88.0
▸ What https://www.puma.com is doing differently...				

JSON-LD is structured data in your page code that tells search and AI systems who you are and what each page represents. You have some JSON-LD, but it looks thin across the pages we sampled—spreading complete, accurate labels to more key templates usually lifts how confidently you get cited.

llms.txt is an optional but increasingly common helper file: a short list of important URLs and notes so AI tools know where to look first. It doesn't replace good content, but it reduces confusion and helps assistants surface the right pages when people ask about your brand.

Measures whether AI search tools can recognise the brand and use the site's pages as clear, cite-worthy answers. AI systems may struggle to recognise the brand or find clear, cite-worthy answers on important pages.

WHAT IS WORKING WELL

AI citability:

- Homepage HTML title is present and usable for basic metadata checks.
- Open Graph / preview images on at least one URL.
- JSON-LD present on at least one URL (successfully parsed in crawl).

AI search success:

- Sampled pages show no nosnippet / max-snippet:0 on meta or X-Robots-Tag.

Brand/entity visibility:

- Automated scan found 2 likely presence signal(s) for "Nike" across Wikipedia/YouTube/Reddit/LinkedIn.

WHAT NEEDS WORK

AI citability:

- Improve Organization structured data: add verified profile links (sameAs) where confirmed and extend JSON-LD on priority templates so AI can recognize and cite your business.

AI search success:

- Set up measurement for AI-search visit quality, including conversions, engaged sessions, and query-to-page reporting.
- Add or verify dateModified / datePublished on time-sensitive templates.

Brand/entity visibility:

- Strengthen authoritative footprint (Wikipedia when eligible, key socials) where it fits brand strategy.
- Manually verify key third-party brand sources such as reviews, directories, news mentions, and Google Business Profile.

Query coverage & citation footprint:

- Add pages that match common AI query shapes (definitions, comparisons, troubleshooting, pricing, selection).
- Check whether the brand appears in ChatGPT, Perplexity, Google AI Overviews, and Copilot for priority queries.

AI platform readiness

These scores estimate how ready the site is for each major AI search experience in this model. They combine crawl signals, brand visibility checks, structured data, content quality, and discovery signals such as llms.txt and sitemap reachability. Treat them as a guide—confirm with live searches and platform tools where it matters.

**Google AI Overviews**

Confidence: medium — automated crawl and signal proxies; confirm in platform-specific tools.

58 /100

58

Add question-style H2/H3s and direct 1–2 sentence answers—AIO favors clarity beyond raw rank (confirm top-10 SEO in Search Console).

**ChatGPT**

Confidence: medium — automated crawl and signal proxies; confirm in platform-specific tools.

62 /100

62

Grow Reddit and YouTube footprint—high citation share for ChatGPT-style web answers.

**Perplexity**

Confidence: medium — automated crawl and signal proxies; confirm in platform-specific tools.

57 /100

57

Earn authentic Reddit and forum visibility—Perplexity weights community validation heavily.

**Google Gemini**

Confidence: medium — automated crawl and signal proxies; confirm in platform-specific tools.

54 /100

54

Expand Schema.org + sameAs—Gemini uses structured data aggressively for entity understanding.

**Bing Copilot**

Confidence: medium — automated crawl and signal proxies; confirm in platform-specific tools.

70 /100

70

Improve discovery signals (live llms.txt + sitemap reachability)—Copilot/Bing ecosystem hygiene.

**Claude**

Confidence: medium — automated crawl and signal proxies; confirm in platform-specific tools.

60 /100

60

Curate llms.txt toward documentation and source-of-truth URLs (optional signal in the Claude rubric).

robots.txt is a simple instruction file at the root of your site that tells crawlers (including many AI crawlers) which URLs they may fetch. We could retrieve yours—still audit it whenever you launch new sections, and avoid overly broad “block everything” rules that can hide you from AI search results.

The **technical checks in this report** look at whether **robots.txt** was available, whether sampled pages returned a normal **successful response** (we report this as an HTTP status such as 200—meaning the page answered OK), and whether our crawl used a **trusted HTTPS** connection. **Core Web Vitals** (real-user loading and layout stability) and **time to first byte** (how quickly the server starts responding) matter for visitors and search quality, but they are *not measured in this automated run*—use your analytics toolkit or a dedicated speed test for those.

Measures whether search engines and AI tools can access the site, find the right pages, and read the important content reliably. The technical foundation is solid, but a few access, discovery, performance, or crawler-readability issues still need attention.

WHAT IS WORKING WELL

Indexability, canonicalisation & crawl health:

- robots.txt fetched successfully.
- Most sampled pages return HTTP 200 (40/40).
- High share of HTTP 200 in sample—good baseline for URL resolution checks.

AI crawler access / robots:

- Tier 1 AI search/retrieval crawlers effectively allowed on hero URL (5/5).
- Googlebot and Bingbot effectively allowed (2/2).
- No blanket `\*` site block and no sampled `noai` signals (+10).

SSR & raw HTML completeness:

- Homepage metadata was readable in raw HTML (good sign for SSR/HTML completeness).
- Crawl could extract some structured/preview signals from HTML.

Performance, mobile, HTTPS:

- TLS verification mode suitable for reliable crawl.
- Most sampled pages returned normally (good baseline for crawl).

llms.txt, sitemaps, freshness:

- Sitemap reachable and allowed for Tier-1 check.

WHAT NEEDS WORK

Indexability, canonicalisation & crawl health:

- Ask your developer to update the site’s crawler access rules in robots.txt so search engines and AI tools can read the public pages you want cited.
- Check key templates to make sure important pages are listable in search and point to the correct main URL.
- Sampled 200 responses include noindex (meta or X-Robots-Tag)—confirm intent; unintentional noindex breaks consolidation.

AI crawler access / robots:

- Decide whether to allow ByteDance’s AI training crawler, Bytespider. This is a data-use policy choice, not a core AI search visibility fix.

Performance, mobile, HTTPS:

- Run a page speed and usability check to confirm whether loading speed is affecting users or crawlers.

llms.txt, sitemaps, freshness:

- Publish a short AI guide file (llms.txt) that points AI tools to your most important product, service, support, company, and policy pages.
- Freshness signals (lastmod, feeds, IndexNow, Last-Modified) are manual checks in this version.

Overall score cap:

- Score cap: 55 — Some sampled pages appear to be marked noindex (blocks indexing/citation).

AI crawler access

This table shows which search and AI crawlers can access the site’s public pages under the current rules. For AI visibility, the practical goal is to let answer-oriented crawlers read the pages you want cited, while deciding separately whether training-only crawlers fit your data-use policy.

► Technical scoring detail (how the composite is built)

CRAWLER	TIER	GEO RECOMMENDATION	REASON	YOUR ROBOTS (HOMEPAGE)	ALIGNED
OAI–SearchBot	1	ALLOW	ChatGPT Search / search retrieval	Can fetch	Yes
ChatGPT–User	1	ALLOW	User-initiated browsing	Can fetch	Yes
GPTBot	1	ALLOW	OpenAI crawler; policy-sensitive visibility	Can fetch	Yes
ClaudeBot	1	ALLOW	Claude web access / retrieval	Can fetch	Yes
PerplexityBot	1	ALLOW	Perplexity search and citations	Can fetch	Yes
Googlebot	2	ALLOW	Google Search; foundational for AI Overviews	Can fetch	Yes
Bingbot	2	ALLOW	Bing index; Copilot / Microsoft search	Can fetch	Yes
Google–Extended	2	ALLOW	Gemini apps / Vertex-style use; not normal Search indexing	Can fetch	Yes
GoogleOther	2	ALLOW	Google non-search product/research crawls	Can fetch	Yes
Applebot	2	ALLOW	Siri / Spotlight / Apple search features	Can fetch	Yes
Applebot–Extended	2	ALLOW	Apple Intelligence extension; policy-sensitive	Can fetch	Yes
Amazonbot	2	ALLOW	Alexa / Amazon AI discovery	Can fetch	Yes
FacebookBot	2	ALLOW	Meta AI-related crawling	Can fetch	Yes
CCBot	3	CONTEXT	Common Crawl / training datasets	Can fetch	—
anthropic–ai	3	CONTEXT	Anthropic training/research (not ClaudeBot)	Can fetch	—
cohere–ai	3	CONTEXT	Cohere model training	Can fetch	—
Bytespider	3	BLOCK	ByteDance; market/abuse considerations	Can fetch	No
meta–externalagent	3	CONTEXT	Meta external AI data collection	Can fetch	—

Content quality & structure

59.5

Weak

(30%)

Measures whether the site's content is helpful, trustworthy, well organised, and easy for AI systems to understand. Important content is not yet clear, trusted, or structured enough for strong AI citation performance.

WHAT IS WORKING WELL

E-E-A-T & helpful content:

Homepage HTML title is present and usable for basic metadata checks.

Open Graph / preview images on at least one URL.

robots.txt fetched successfully.

Most sampled pages return HTTP 200 (40/40).

Passage-level answerability:

JSON-LD present on at least one URL (successfully parsed in crawl).

WHAT NEEDS WORK

Original information gain:

Add original information gain (first-party data, case studies, expert commentary, unique examples).

Passage-level answerability:

Improve Organization structured data: add verified profile links (‘sameAs’) where confirmed and extend JSON-LD on priority templates so AI can recognize and cite your business.

Add passage patterns AI can lift: definitions, direct answers under headings, tables, steps, FAQs, key takeaways.

Source transparency & governance:

Improve source transparency (named sources, dates, methods) and content governance (review process, corrections, disclosures).

Trust signals from your crawl (E-E-A-T view)

Four cards translate Experience, Expertise, Authoritativeness, and Trust into plain language. Each score is an automated blend of crawl checks—not a substitute for editorial, legal, or reputation review.

E-E-A-T stands for **E**xperience, **E**xpertise, **A**uthoritativeness, and **T**rustworthiness—how well your content demonstrates real-world experience, subject-matter depth, credible sourcing, and reliability. Search and AI systems use these signals to decide what to surface and cite. The scores below are *approximate indicators* from this crawl; they complement—not replace—a full editorial and fact-check review.

Experience

53/100

Does the content feel grounded in real use or practice?

Experience is about whether pages read like they were written by people who have actually done the job, used the product, or lived the situation—not just generic filler. Search and AI systems reward depth, specificity, and helpful detail that sounds first-hand.

How we calculate this number:

We take 55% of your overall content quality & structure score from this audit (that score already mixes several automated content checks). We then add up to 10 points when sampled pages had social preview images, and up to 10 more when structured data appeared—both tend to travel with richer, more complete pages. This crawl applied: share/preview images on sampled pages (+10); structured data snippets (+10).

Expertise

42/100

Does the site show clear subject-matter depth?

Expertise is how convincingly you demonstrate knowledge: clear explanations, useful structure, and cues that a knowledgeable author or organization stands behind the content. AI cites sources that look like they know the topic end-to-end.

How we calculate this number:

We take 45% of your content quality & structure score, then add 15 points when structured data was detected (it helps machines understand entities and facts) and 5 points when a live llms.txt file exists (a small discoverability signal for AI tools). This crawl applied: structured data present (+15).

Authoritativeness

31/100

Can others (and machines) verify who you are?

Authoritativeness is recognition and corroboration: consistent branding, visible footprint, and links that tie your site to trusted profiles elsewhere. When those line up, readers and models treat you as a real entity worth citing.

How we calculate this number:

We take 60% of your brand / entity visibility score from this audit, then add 15 points when structured data includes sameAs links to official profiles (for example social or knowledge-base URLs you control). This crawl did not find sameAs profile links in structured data, so only the brand-visibility blend counted.

Trust

78/100

Do the basics feel safe, maintained, and accessible?

Trust covers the hygiene people notice before they read a word: secure browsing, pages that load without errors, and sensible rules for bots. It is not a legal or financial audit—just whether automated checks saw a well-maintained, reachable site.

How we calculate this number:

We blend 55% of your technical setup score with 25% of your AI crawler access score, then add 10 points when HTTPS/TLS looked healthy in this crawl and 5 points when llms.txt was present. This crawl also applied: HTTPS/TLS looked healthy (+10).

Brand visibility

This section checks whether the brand **Nike** is visible on major third-party platforms that AI systems often use for verification. Automated matches should be confirmed manually before treating a URL as an official profile or adding it to structured data.

Checks: Wikipedia (multi-query API), YouTube (@handles + search HTML), Reddit (search.json for up to three on-brand threads; optional DistilBERT SST-2 sentiment on title + top-scoring comments when requirements-brand-sentiment.txt deps are installed; else subreddit slug fallback), LinkedIn (expanded /company/ slugs). Pass an accurate --brand and use the audited domain so host-based guesses (e.g. eurocarparts → Euro Car Parts) apply. Confirm in browser if a site blocks bots.

PLATFORM	PRESENCE	CHANNEL / PAGE	IMPACT ON AI VISIBILITY
Wikipedia	Likely yes	Article: Nike, Inc. Open	Very high — frequent citation source for ChatGPT and other models.
YouTube	No / unclear	No channel matched handles or search results; verify manually.	Medium–high — rich media and entity linkage.
Reddit	No / unclear	No Reddit search hits for brand variants; try site:reddit.com in a browser.	Medium — community discussions models may weight.
LinkedIn	Likely yes	Company page resolves (slug try: nike). Open	High — company and executive verification signals.

Produced by create-report.py. Scores are heuristic; validate with live checks and analytics.

Sample scripts

Suggested robots.txt merge, json-ld sample, and llms.txt skeleton from this run. Expand to scroll and copy.

robots.txt (merged suggestion)

Expand · scroll inside · copy

json-ld.txt (WebSite sample)

Expand · scroll inside · copy

llms.txt (generated)

Expand · scroll inside · copy

